Artificial Intelligence
3rd year, 1st semester

State based models for decision problems

Week 2: models and uninformed search strategies

"Intelligence is what you use when you don't know what to do" Jean Piaget

The AI expectation on acting rationally:

- we define a problem
- the computer solves it
- the computer outputs the result (unexplainable AI) or the solution (explainable AI)

First approach: state-based models

● in what way should I change the current *state* of the problem in order to get closer or reach the goal?
● compute a solution (algorithm - sequence of steps) starting from an initial *state* and ending in the goal *state*
● mostly covered by search strategies, reasoning systems, AI for games
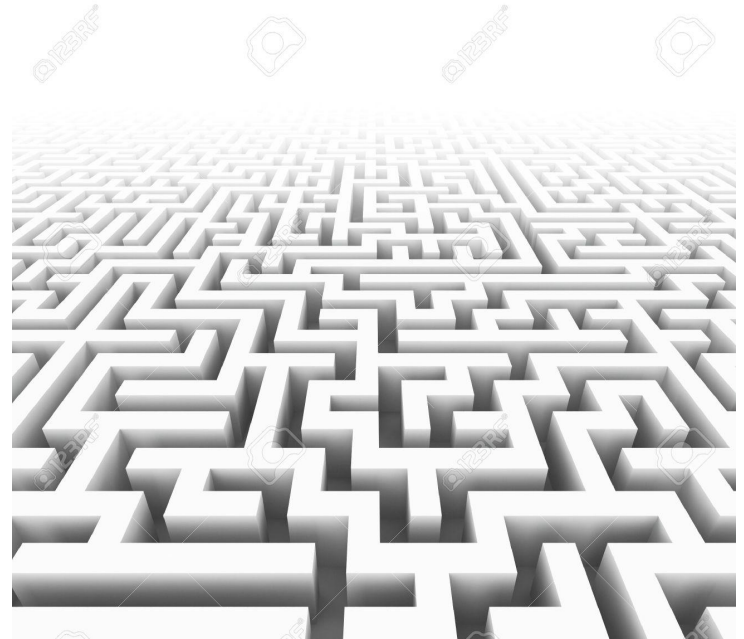
Solving problems

Problem: We have two numbers, A and B. Which is larger?

Instance: A=?, B=?

A solution should produce the correct answer for **all** possible instances

What's the Problem with problems?

P = set of problems that can be solved/checked in Polynomial time by a deterministic TM.

NP = set of problems that can be solved in Polynomial time by a Non-deterministic TM. Checking a NP solution is a P problem.
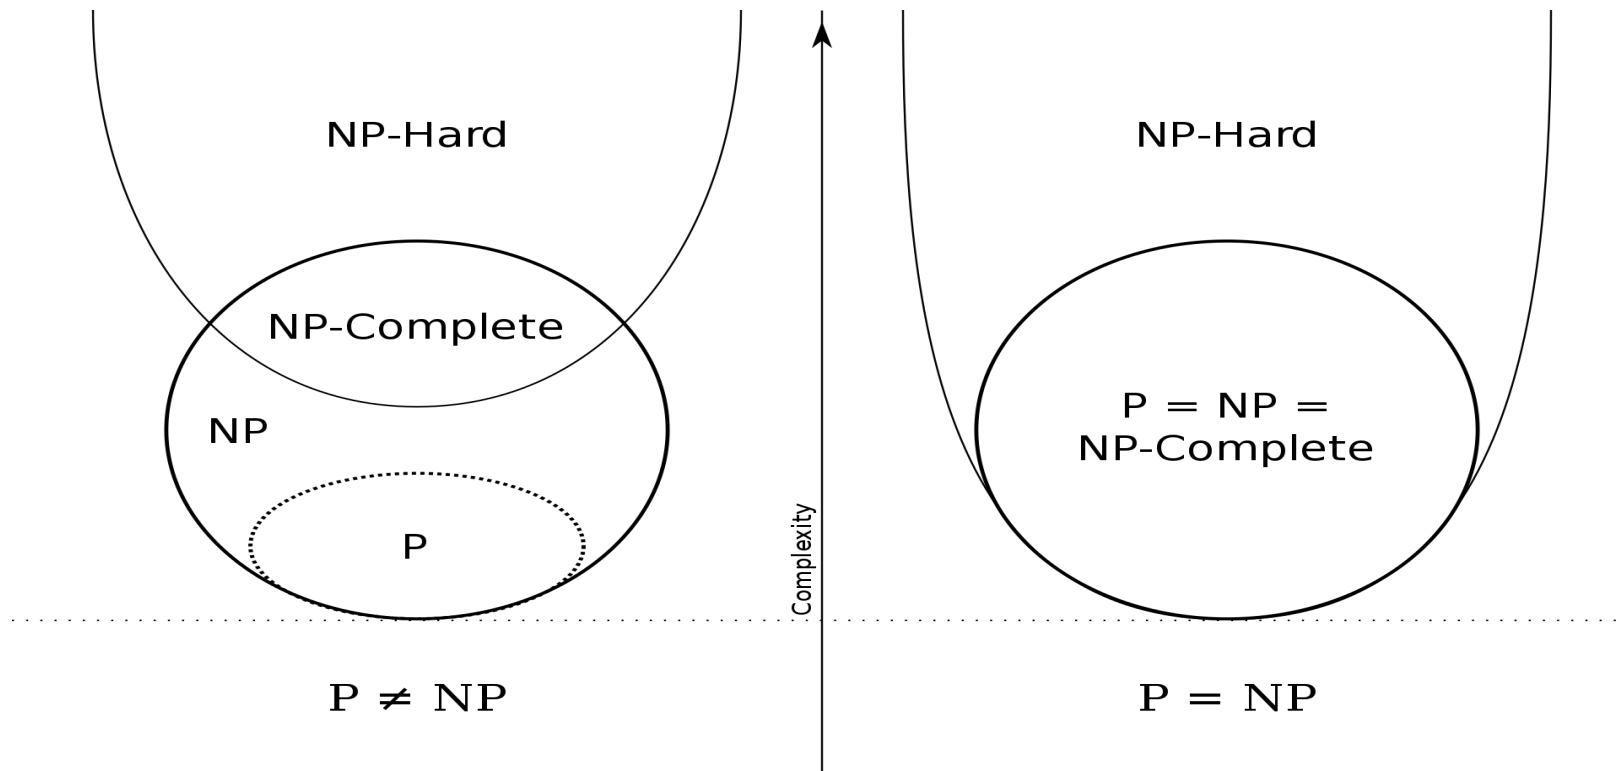
NP-complete = subset of NP with the property that any problem in NP can be reduced to a NP-complete problem in polynomial time.

NP-hard = set of problems with the property that any problem in NP can be reduced to a NP-hard problem in polynomial time.

P=NP ?

NP-complete problems

Courtesy of Wikipedia

Let the computer solve our problems! Or not...

Solvable NP-hard problems: QSAT: is a logical proposition using quantified (existential $\exists$ or universal $\forall$) variables satisfiable?

Unsolvable NP-hard problems: Turing Halting problem: given an algorithm and an input, determine whether the algorithm execution for that input will eventually halt.

The rest can be solved by a computer.

Don't think about solutions, think about problems.

Find solutions requires effort, describing problems requires intelligence.

"If I had only one hour to save the world, I would spend fifty-five minutes defining the problem, and only five minutes finding the solution." A.Einstein

Let's find a NP-complete problem and a computer able to solve it

Problem: finding a path in a graph

Describing a model - reducing a problem

- describe a state
- identify special states and the problem space
- describe the transitions and validate them
- specify a search strategy

Choosing a representation for a state

State: all data required to continue looking for a solution

No ambiguity

Expressive enough to include all required data

Choosing a representation for a state

Compact enough to be easy to store and perform changes on it

A list of towers pieces are placed, in the order of piece size size

$$(3, 3, 3, 3, 1, 1, 2, 2)$$

Is this good enough?

Choosing a representation for a state

Compact enough to be easy to store and perform changes on it

A list of towers where pieces are placed, in the order of piece size, with the number of towers added at the beginning

$$(\textcolor{red}{3}, 3, 3, 3, 3, 1, 1, 2, 2)$$

$$(n, t_1, t_2, ..., t_m), \ 1 \leq t_i \leq n$$

Choosing a representation for a state is the most **important and difficult** step

No ambiguity

Compact

Expressive

Includes all required data

There is no generally best representation

Special states

Initial state (3, 1, 1, 1, 1, 1, 1, 1, 1)

State Initialize (int n, int m)
{
      Return (n, 1, 1, 1, …, 1);
}

m

There can be more than one initial state. There has to be at least one initial state.

Special states

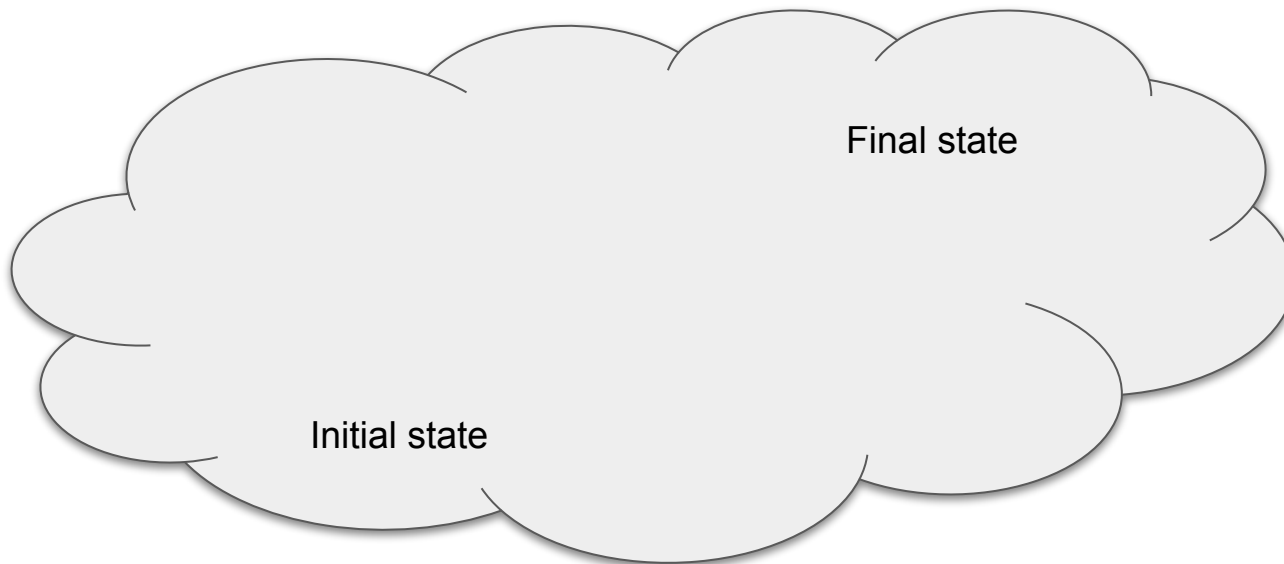Final state(s) (n, n, n, n, n, n, n, n, n)

Boolean IsFinal (State s)
{
    If s = (k, k, k, …, k) then return true;

          m+1

    else return false;
}

There can be more than one final state. There has to be at least one final state.

FACULTY OF
COMPUTER SCIENCE

IA@FII

Problem space: all possible states

Final state

Initial state

Finite space (limited variability in chosen state representation and only valid states are included)
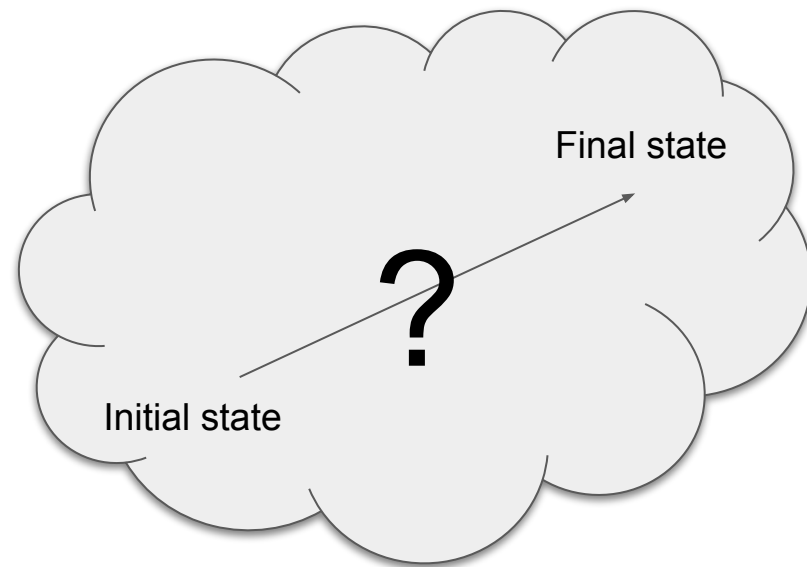
What is the problem now?

How difficult is it to solve?

Problem space has *m* dimensions

How many states are within the space?

How can you move within the
 problem space?

Simon's Ant

Complexity of the solution is given by the problem, NOT by the way in which you solve it

Transitions

Only one possible way to change current state: move one piece to another tower.

$(n, t_{11}, t_{12},...,t_{1m}) \rightarrow (n, t_{21}, t_{22},...,t_{2m})$, where $t_{1i} = t_{2i}$ for all $1 \le i \le m$, except exactly one $i = k$

State Transition (State s, piece, tower)

Valid transitions

1. No smaller piece is placed atop piece k
2. No smaller piece ends up below piece k

1. $t_{1i} \neq t_{1k}$, for all $1 \leq i < k$
2. $t_{2i} \neq t_{2k}$, for all $1 \leq i < k$

Boolean Validate (State s, piece, tower)

Implement transitions and validation separately, it's good practice!

Search strategy

```
Void strategy(State s)
{
    While (!isFinal(s))
        {
            Choose piece, tower;
            If (Validate (s, piece, tower))
                s = Transition(s, piece, tower);
        }
}
```

Search strategies
- Uninformed - no distinction between states
  - Random
  - BFS and Uniform Cost
  - DFS and Iterative Deepening
  - Backtracking
  - Bidirectional
- Informed - heuristics to help distinguish between states
  - Greedy best-first
  - Hillclimbing
  - Beam
  - A* and IDA*

Breadth First Search and Uniform Cost

- Visit states in the order of distance (number of transitions) from the initial state.
- Explores all immediate neighbors (accessible states) until no more neighbors or final state found.
- Has to memorise each generated state: very costly.
- Might visit the same state multiple times.
- Finds shortest path (optimum solution)
- Uniform cost: if options have different costs, explore cheaper paths first
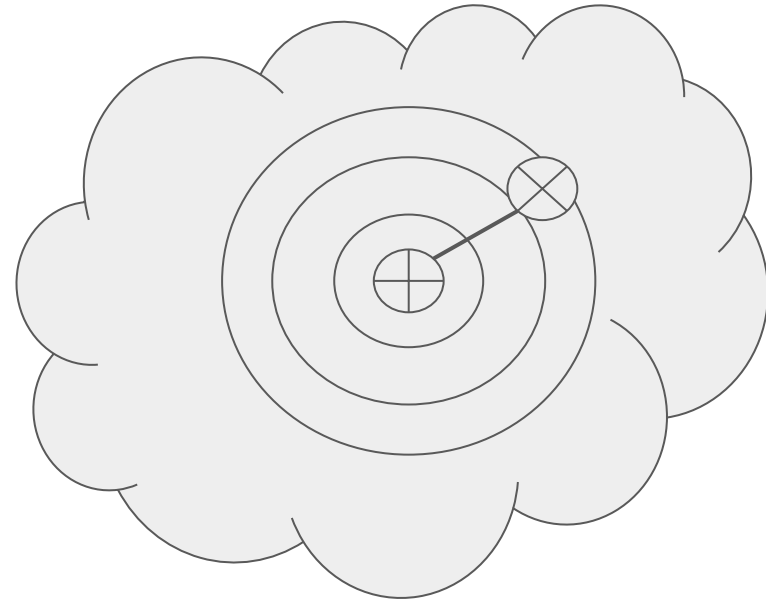
Depth First Search and Iterative Deepening Search

- Visit one immediate neighbor of the current state until final state is found, return to previous unexplored neighbor if no more neighbours available for current state.
- Has to memorise each generated state: very costly.
- Might visit the same state multiple times.
- Might not finish (if loops are present in the problem space)
- IDS: Explore only up to an increasing depth (distance from initial state) - no more infinite paths
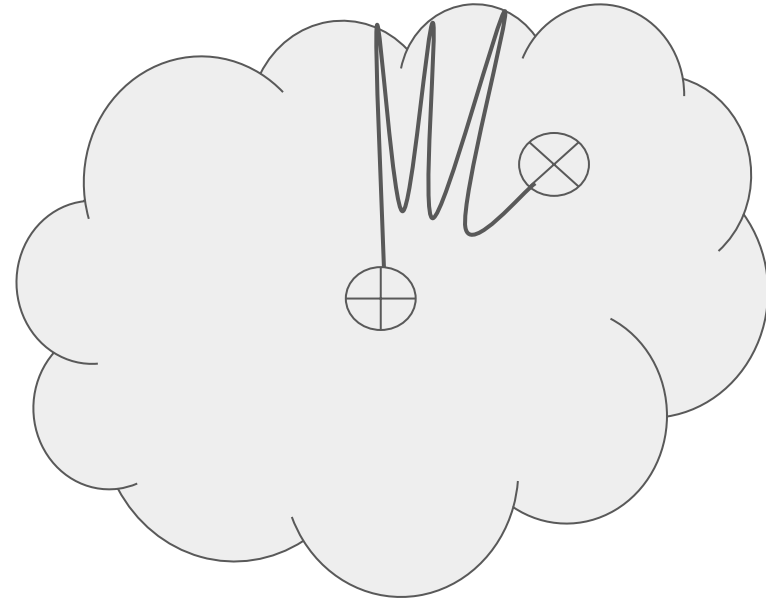
Random vs DFS
- The same, if random remembers visited neighbours

Backtracking vs DFS
- NOT the same
- Backtracking sets an order for the neighbor states
- How do you order states?
- Backtracking doesn't need to memorize visited states!
- Backtracking can avoid loops WITHOUT storing visited states

Bidirectional Search

- Starts exploring from both the initial and final state simultaneously
- Has to memorise each generated state, but they should not be that many
- Finds shortest path (optimum solution)
- What are the reverse transitions?
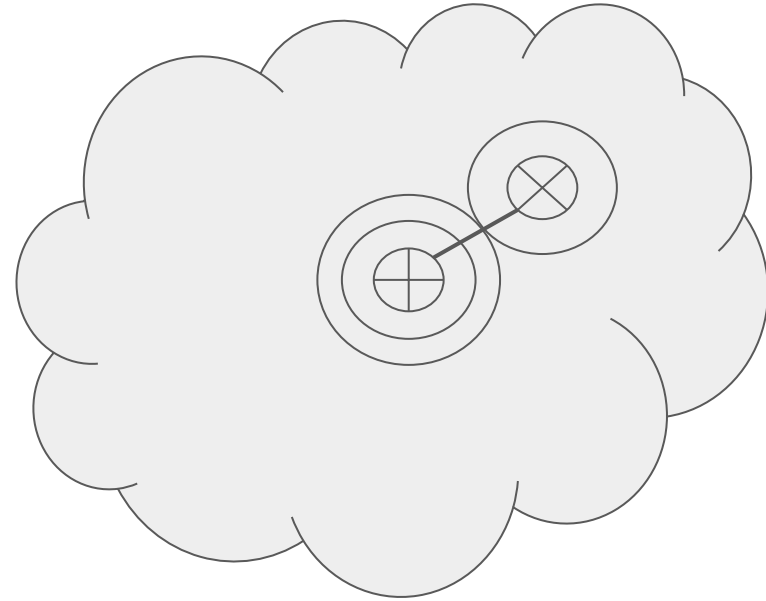
Comparison

Complexity of the solution is given by the problem, NOT by the way in which you solve it

Criterion	BFS	DFS	IDS	Bidirectional	BKT	Random
ETC	N^O	N^A	N^O	N^O	N^A	N^A
Optimum	Yes	No	Yes	Yes	No	No
All	Yes	No	Yes	Yes	Yes	No

N = average number of accessible states
O = length of the optimum solution
A = length of the average solution

References (for search strategies)

https://w3cschoool.com/uninformed-search-algorithm

Stanford Programming Abstractions:
https://www.youtube.com/playlist?list=PLFE6E58F856038C69